

Introduction to Data Science

dr Maciej Świtała
Ewa Weychert

Class 8: AI and prompt engineering

e.weychert@uw.edu.pl

What is AI?

- **Artificial Intelligence (AI)**: systems that perform tasks that normally require human intelligence.
- Examples:
 - Understanding and generating text.
 - Recognizing images, audio, and video.
 - Making predictions and recommendations.
- Today we focus on **Large Language Models (LLMs)** like ChatGPT.

Claim: “There is no AI, only LLMs”

- You sometimes hear: There is no AI, only LLMs.
- This statement feels true because:
 - Most visible AI apps today = chatbots based on LLMs.
 - People interact daily with ChatGPT, Claude, Gemini, etc.
- But from a scientific perspective, this is **not accurate**.
- Let’s unpack what **AI** and **LLMs** actually are.

What Is AI? (Broad Definition)

- **Artificial Intelligence (AI):**
 - Any technique that makes computers perform tasks which, if done by humans, would require intelligence.
 - AI includes many subfields:
 - Classical search, planning, logic.
 - Expert systems, rule-based systems.
 - Machine learning, deep learning.
 - Computer vision, speech recognition, robotics.
 - **Large language models (LLMs).**
 - **Conclusion:** LLMs are *one branch* of AI, not the whole tree.

What Is an LLM?

- **Large Language Model (LLM):**
 - A model trained to predict the **next token**
 - in a sequence of text, given the previous tokens.
 - Trained on massive text corpora:
 - Books, websites, code, documents, etc.
 - Core capability:
 - Generate and transform text (and code) by pattern matching.
 - They do **not** have:
 - Human-like understanding or consciousness.
 - Persistent memory (unless wrapped in a larger system).

Why It Feels Like “AI = LLMs” Today

- Most popular AI tools are LLM-based:
 - ChatGPT, Claude, Gemini, Copilot, etc.
- They are:
 - General-purpose (text, code, reasoning, explanation).
 - Accessible through a simple chat interface.
 - Capable of talking like a human.
- For many non-technical users:
 - AI = the chatbot I talk to.
- This drives the slogan:

There is no AI, only LLMs.

Why That Claim Is Misleading

- It ignores the rest of AI:
 - Vision, robotics, control, planning, classical ML, etc.
- It confuses:
 - **Models** (LLMs)
 - with full **systems** (AI agents, copilots, apps)
- It encourages **anthropomorphism**:
 - Treating a text predictor as a real intelligence.
- Better statement:

Most visible AI apps today are built around LLMs.

LLMs as Components, Not Whole Systems

- Modern AI assistants often combine:
 - An LLM (language + reasoning engine).
 - Tools (search, code execution, databases, APIs).
 - Memory (storing user context, documents).
 - Orchestration / planning logic.
- The LLM is:
 - The **interface** (natural language).
 - One part of a larger **AI system**.
- Saying only LLMs hides all other components.

LLMs Are Powerful, But Not AGI

- LLMs can:
 - Simulate reasoning.
 - Explain concepts.
 - Generate complex artifacts (code, essays, plans).
- However, they:
 - Do not have goals or desires.
 - Do not understand the physical world.
 - Do not build internal causal models like humans.
- Important for students:
 - Respect the power of LLMs.
 - But avoid calling them **full intelligence**.

A More Accurate Way to Say It

- Instead of:

There is no AI, only LLMs.

- Use:

Most of what people call AI today are systems built around large language models.

- Or:

Modern AI applications often use LLMs as a core component, but AI as a field is much broader.

Key Takeaways for Students

- **AI** is a broad field; **LLMs** are one powerful technique within it.
- LLMs:
 - Are next-token predictors trained on huge text corpora.
 - Do not equal general intelligence.
- Most popular AI apps today are LLM-centric, which fuels the confusion.
- As ML/DS practitioners:
 - Understand LLMs deeply.
 - But also know the wider AI toolbox beyond LLMs.

Why Prompts Matter

- The model **does not read your mind**—it only sees your text.
- Small changes in wording can lead to very different outputs.
- Good prompting turns a generic model into:
 - A tutor.
 - A software architect.
 - A marketing expert.
 - A research assistant.
- Prompting is a **skill** you can improve with practice.

How LLMs Work (Intuition)

- LLMs are trained to **predict the next token** (piece of text).
- They do not have:
 - Human-like understanding.
 - Persistent memory of you (unless explicitly provided).
- They **simulate** patterns of reasoning based on training data.
- Your prompt defines the **context** and **goal**.

Prompting

What is a Prompt?

- A **prompt** is any input you give the model:
 - Instructions (Summarize this article).
 - Data (a paragraph, table, code snippet).
 - Examples (good / bad outputs).
- Prompts can be:
 - **Unstructured:** natural language instructions.
 - **Structured:** clear sections, headings, bullet points, formats.

Anatomy of a Good Prompt

A useful mental checklist:

- **Role** — Who is the AI?
- **Task** — What should it do?
- **Context** — Background, constraints, domain.
- **Output Format** — Bullet list, table, JSON, etc.
- **Tone & Audience** — Level of detail, style, jargon.

Formula:

You are [role]. Your task is to [task]. Consider this context: [context]. Produce output in [format] for [audience].

Clarity & Specificity

- Avoid vague prompts:
 - Explain AI.
- Add constraints and details:
 - Length (150 words).
 - Audience (non-technical managers).
 - Purpose (so they can decide whether to invest).

Example

Instead of:

Explain AI.

Use:

In about 150 words, explain what AI is to non-technical managers so they can decide whether to invest in it. Use simple language and a concrete example.

Bad vs Good Prompts

Weak prompt

Write about dogs.

Problems:

- No audience.
- No purpose.
- No length or style.

Improved prompt

Write a 150-word, beginner-friendly explanation of how guide dogs are trained. Use simple language and concrete examples so a 12-year-old can understand.

Improvements:

- Clear topic.
- Defined audience.
- Length + style.

Using Constraints Effectively

- **Length:** word count, paragraphs, bullet points.
- **Tone:** formal, friendly, humorous, neutral.
- **Perspective:** first person, third person, etc.
- **Structure:** sections, headings, numbered steps.
- **Language:** specify language, level (A2, B2, etc.).

Example

Create a one-page executive summary (about 300 words) in bullet points and sub-bullets that a busy CEO can read in 2 minutes.

Examples: Few-Shot Prompting

- Provide **examples** of what you want.
- The model learns your pattern within the conversation.

Example

Prompt:

You are a writing assistant. Rewrite sentences to be clearer and more concise.

Example 1:

Input: Due to the fact that we were late, we missed the train.

Output: Because we were late, we missed the train.

Now improve this sentence:

Input: In light of recent events, it has become increasingly necessary to improve our processes.

Bad vs Good Prompts: Data Cleaning (Missing Values)

Weak prompt

Clean this dataset.

Problems:

- No information about columns.
- No definition of clean.
- No preferred tools or language.

Improved prompt

You are a data analyst. I have a CSV file with columns: `age`, `income`, `city`, `joined_date`.

1. Identify which columns have missing values and how many.
2. Propose sensible strategies to handle the missing values in each column (e.g. drop, mean/median, category Unknown).
3. Give me Python `pandas` code that:
 - prints the missing value counts,
 - applies your chosen strategies,
 - returns a cleaned DataFrame called `df_clean`.

Improvements:

- Clear columns and context.

Bad vs Good Prompts: Data Types & Dates

Weak prompt

Fix the types in my data.

Problems:

- No column names.
- No target types.
- No expected output format.

Improvements:

- Clear mapping from columns to target types.
- Request for both explanation and code.
- Explicit assumptions.

Improved prompt

You are a data engineer working in Python. I have a `pandas` DataFrame `df` with columns:

- `user_id` (should be string),
 - `age` (should be integer),
 - `signup_date` (string in format `YYYY-MM-DD`),
 - `is_active` (values yes/no).
1. Explain which data types are appropriate for each column.
 2. Write `pandas` code that converts each column to the correct type, parses `signup_date` as `datetime`, and converts `is_active` to boolean.
 3. Briefly explain any assumptions you make.

Bad vs Good Prompts: Outliers

Weak prompt

Find outliers in my data.

Problems:

- No definition of outlier.
- No columns specified.
- No request for method transparency.

Improved prompt

Act as a data scientist. I have a DataFrame `df` with numeric columns `price` and `quantity`.

1. Suggest at least two methods to detect outliers in these columns (e.g. z-score, IQR).
2. For each method, explain in 2–3 sentences how it works and when it might fail.
3. Provide Python code using `pandas` and `numpy` that flags outliers in `price` and `quantity` using the IQR rule ($1.5 * \text{IQR}$).
4. Return the code in a single code block.

Improvements:

- Clear columns and expectations.
- Multiple methods + short explanations.
- Ready-to-run code requested.

Role Prompting

- Assign the model a **role** to shape style & content:
 - You are a senior data analyst.
 - You are a patient math tutor.
 - You are an experienced UX researcher.
- Roles influence:
 - Vocabulary and tone.
 - Depth of detail.
 - Perspective and focus.

Example

You are a senior product manager. Explain to a junior colleague why clear requirements are critical for successful software projects. Use practical examples.

Chain-of-Thought Prompting

- Ask the model to **think step by step**.
- Helps for reasoning, planning, debugging, learning.
- Useful phrases:
 - Think aloud as you solve this.
 - Explain your reasoning step by step.
 - First outline the approach, then give the answer.

Example

Show me step by step how you would break down this problem before giving the final answer: [problem].

Decomposing Complex Tasks

- Break large tasks into **smaller stages**:
 - ① Understand the problem.
 - ② Generate options.
 - ③ Evaluate pros/cons.
 - ④ Produce final output.
- Run the model multiple times, one step at a time.

Example

We will work in stages.

Step 1: Ask me up to 5 clarification questions about my project.

Step 2: Propose 3 different solution approaches.

Step 3: After I choose one, create a detailed plan.

Iterative Refinement

- Treat AI as a **collaborator**, not a vending machine.
- Cycle:
 - ① Draft.
 - ② Critique.
 - ③ Revise.
- Use prompts like:
 - Critique this and suggest 3 improvements.
 - Rewrite this with your suggestions applied.

Example

Here is my email draft. First, give me a brief critique (max 5 bullet points). Then rewrite the email applying your suggestions.

Prompt Templates

- Create **reusable templates** for common tasks:
 - Emails.
 - Meeting summaries.
 - Code reviews.
 - Report outlines.
- Store them in a note or document for quick copy–paste.

Template Example: Email Draft

You are a professional communication coach. Draft a clear, polite email to [recipient role] about [topic]. Include: (1) short context, (2) main request, (3) next steps. Use a friendly but concise tone.

System vs User Instructions

- Many tools support a higher-level **system** message:
 - Global behavior, rules, persona.
- **User** message:
 - Specific task or question.
- Keep system instructions:
 - Stable, high-level.
 - Clear about boundaries & style.

Example System Instruction

You are a careful assistant for software engineers. Always ask 1–2 clarification questions when requirements are ambiguous. Prefer concise, practical examples in code.

Structured Outputs

- Ask for output in a **specific structure**:
 - Tables.
 - Bullet lists.
 - JSON / YAML.
- Helps when:
 - You want to import into other tools.
 - You need consistent formatting.

Example (JSON)

Summarize this document into JSON with keys: "audience", "main_points" (list of strings), and "action_items" (list of strings). Return only valid JSON.

Working with Tools (Conceptually)

- Modern AI systems can use tools:
 - Code execution.
 - Web browsing.
 - File uploads (e.g., PDFs, CSVs).
- Good prompts clarify:
 - What tool to use (if you can specify).
 - What you expect as a result.
- Example: Analyze this CSV and produce 3 key insights and 2 charts.

Safety, Ethics, & Limitations

- Treat AI outputs as **drafts**, not ground truth.
- Watch for:
 - Factual errors and hallucinations.
 - Bias or unfair assumptions.
 - Leaking confidential information.
- Good practices:
 - Do not paste highly sensitive data.
 - Cross-check important facts.
 - Be transparent when AI is involved.

Demo 1: Improving a Weak Prompt

- Start with a weak prompt, e.g.:
 - Write a project plan.
- Show the mediocre output.
- Then refine the prompt live:
 - Add role (You are a project manager).
 - Add context (Software rollout for 200 employees).
 - Add structure (phases, timeline, risks).
- Compare outputs and discuss.

Demo 2: Building a Template

- Choose a task your audience cares about:
 - Meeting notes.
 - User interview summaries.
 - Bug reports.
- Live-build a prompt template:
 - Ask the room: What sections do we need?
 - Turn that into a reusable prompt.
- Save the final template as a prompt asset.

Demo 3: Multi-Step Workflow

- Show a small workflow, e.g.:
 - ① Research topic ideas.
 - ② Select one and outline.
 - ③ Draft a blog post.
 - ④ Generate a summary and social media post.
- Emphasize:
 - Iteration and feedback.
 - Reusing context.
 - Controlling tone and audience.

Business Use Cases

- **Marketing**

- Campaign ideas.
- Copy variants for testing.

- **Product / UX**

- User story drafts.
- Interview guides and summaries.

- **Operations**

- Process documentation.
- Checklists and SOPs.

- **Data / Analytics**

- Exploratory analysis.
- Hypothesis generation.

Personal Use Cases

- **Learning**

- Explain topics at different levels.
- Create quizzes and flashcards.

- **Writing**

- Drafts, outlines, rewrites.
- Grammar and clarity improvements.

- **Planning**

- Travel itineraries.
- Study plans and schedules.

- **Idea generation**

- Project ideas.
- Hobbies and side projects.

Key Takeaways

- Clear, specific prompts **significantly** improve results.
- Use the **Role–Task–Context–Format–Audience** checklist.
- Break big problems into **multi-step workflows**.
- Treat AI outputs as **drafts**, and iterate.
- Build your own library of **prompt templates**.

Next Steps

- Pick one task you do every week.
- Design a prompt template for it.
- Use it for 2 weeks and refine.
- Share good prompts with your team.

Next Steps

- <https://medium.com/@pvnsripati/as-a-developer-struggling-with-ai-prompts-master-ai-prompt-engineering-now-0616c8e74165>
- <https://pbiecek.github.io/books/>
- <https://github.com/pbiecek/ema>

Q&A

Questions?