

INTRODUCTION TO DATA SCIENCE

Introduction to statistics and statistical hypothesis testing

Maciej Świątała, PhD

Autumn 2025



UNIVERSITY
OF WARSAW



FACULTY OF
ECONOMIC SCIENCES

1 Introduction to statistics

- Basic concepts
- Data sources, data collection methods
- Data cleaning and validation
- Population vs. sample, sampling methods
- Random variables and their distributions
- The most common statistical distributions

2 Statistical hypothesis testing

- Type I and type II errors
- The hypothesis testing procedure
- Example: one-sample t-test
- Example: two-sample t-test
- Example: two-sample t-test (left- and right-tailed)
- What is the p-value?

Introduction to statistics

What is statistics?

Statistics, in its intuitive sense, is a **set of methods** that allow us to draw conclusions from data - that is, to transform **raw data** into **knowledge** and ultimately use that knowledge to make **decisions**.

Stages of a typical analytical project

- ① Identify the research or business problem.
- ② Formulate the project objective.
- ③ Collect data.
- ④ Clean and validate data.
- ⑤ Perform exploratory data analysis.
- ⑥ Model building and validation.
- ⑦ Inference.

Research vs. business problem

Research problem:

- an issue that motivates conducting a study,
- should concern phenomena of significant importance - whose better understanding supports socio-economic development,
- facilitates understanding of social issues and supports decision-making and conflict resolution.

Business problem:

- a situation, challenge, or difficulty that limits the efficiency, profitability, or growth of an organization and requires action to solve it,
- has a practical dimension - it is not about discovering new knowledge but about improving organizational performance, financial results, service quality, customer satisfaction, operational efficiency, or competitiveness.

Project objective; research questions vs. hypotheses

- The **project objective** specifies how the analyst intends to address the research or business problem - it identifies the purpose of the analysis and what should result from it.
- Research questions and hypotheses narrow the scope of inquiry:
 - a **research question** is a clear, focused, specific, and usually open-ended question that the researcher aims to answer to gain insight or understanding of a phenomenon.
 - a **research hypothesis** is a detailed, testable statement predicting the outcome of a study, based on existing theory or prior research.
- Typically, research questions are formulated in qualitative studies (based on non-numeric data such as surveys, interviews, focus groups, observations), while research hypotheses are used in quantitative studies (based on numeric data).

Data sources

Primary data (collected directly):

- surveys and questionnaires,
- individual interviews,
- focus groups,
- direct or participatory observation,
- laboratory or field experiments,
- psychometric tests and measurements.

Secondary data (existing sources):

- corporate databases (CRM, ERP, sales logs),
- public/government data (statistics, demographics),
- scientific and industry publications,
- social media and web data,
- data from partners or suppliers,
- historical or archival data.

Data collection methods

Quantitative methods:

- standardized surveys,
- controlled experiments,
- analysis of transactions and system logs,
- KPI monitoring,
- technical tests and measurements,
- GPS data analysis.

Qualitative methods:

- in-depth interviews,
- participant observation,
- case studies,
- content analysis (texts, media),
- ethnographic and field research.

Channels and technologies for data collection

Offline/traditional:

- paper surveys,
- face-to-face interviews,
- field observations.

Online/digital:

- online surveys, mobile forms,
- system logs, app data,
- social listening, media crawling,
- external APIs.

Automatic sensors:

- IoT, RFID, GPS,
- image recognition cameras,
- machine and production monitoring.

Types of data structures

- ① **Cross-sectional data** - many entities observed once, e.g.
 - household income survey in a given year,
 - customer satisfaction survey after a purchase.
- ② **Time series** - one entity observed multiple times over intervals, e.g.
 - monthly revenues of a company from 2015–2025,
 - daily stock price of a listed company.
- ③ **Panel data** - many entities observed repeatedly over time, e.g.
 - annual sales data for 50 firms from 2010–2020,
 - exam results of students across schools for five years.

Data cleaning and validation

Data cleaning - the process of detecting and correcting (or removing) inaccurate records from a dataset, including:

- detection and handling of missing data,
- identifying and treating outliers,
- removing duplicates,
- correcting data entry errors,
- ensuring consistent data types and units.

Data validation - verifying that the data meets quality and accuracy standards:

- logical consistency (e.g., age cannot be negative),
- range and domain checks,
- referential integrity between tables,
- comparing with external or reference data.

Basic concepts

A **population** is the entire set of individuals, objects, or events possessing some common observable feature, about which we want to draw conclusions.

A **sample** is a subset of the population selected for observation or measurement, usually to make inferences about the entire population.

Sampling methods

Random sampling:

- simple random sampling - each unit has an equal chance of selection,
- systematic sampling - selecting every k -th unit from an ordered list,
- stratified sampling - dividing the population into subgroups (strata) and sampling within each,
- cluster sampling - selecting entire groups or clusters.

Non-random sampling:

- convenience sampling,
- judgmental or expert-based sampling,
- snowball sampling (for hidden populations).

Random variables

A **random variable** assigns a real number to each outcome of a random experiment; it can take discrete or continuous values.

Types of random variables:

- **discrete random variable:** countable number of outcomes, example: number of defective items in a batch.
- **continuous random variable:** takes values from an interval of real numbers, example: body temperature, height, income.

Probability distributions

Probability distribution describes how probabilities are assigned to possible values of a random variable.

We describe distributions with:

- ① **Probability Mass Function (PMF)** (for discrete distributions):

$$\sum_i P(X = x_i) = 1, \quad P(X = x_i) \geq 0$$

- ② **Probability Density Function (PDF)** (for continuous distributions):

$$\int_{-\infty}^{\infty} f(x) dx = 1, \quad f(x) \geq 0$$

Cumulative Distribution Function (CDF)

$$F(x) = P(X \leq x)$$

The most important properties:

- applicable for both discrete and continuous distributions,
- for discrete variables it is a step function,
- for continuous variables: smooth, increasing function,
- $0 \leq F(x) \leq 1$,
- $\lim_{x \rightarrow -\infty} F(x) = 0$, $\lim_{x \rightarrow \infty} F(x) = 1$.

Parameters of a random variable distribution

- Distributions have parameters that determine their final mathematical form.
- Several basic types of parameters are distinguished, which define the distribution's shape, location, and scale:
 - ① **location parameters** – determine the center of the distribution, e.g. the mean, or the shift of the normal distribution, i.e. μ in $N(\mu, \sigma^2)$.
 - ② **scale parameters** – describe the spread of values around the center, e.g. the variance of the normal distribution, i.e. σ^2 in $N(\mu, \sigma^2)$.
 - ③ **shape parameters** – define the general shape of the distribution, its asymmetry, tail heaviness, or peakedness, e.g. p in the Bernoulli or Binomial distribution, or the number of degrees of freedom in the t-Student, Chi-square, and F-Snedecor distributions.
 - ④ **mixed/additional parameters**, e.g. λ in the Poisson or Exponential distribution, describing the rate or intensity of events.

The most common statistical distributions

Discrete distributions:

- Bernoulli,
- Binomial,
- Poisson.

Continuous distributions:

- Uniform,
- Normal (Gaussian),
- Exponential,
- Chi-square,
- Student's t,
- F-Snedecor.

Notation of the most common statistical distributions

Discrete distributions:

- Bernoulli: $X \sim \text{Bern}(p)$,
- Binomial: $X \sim \text{Bin}(n, p)$,
- Poisson: $X \sim \text{Poisson}(\lambda)$.



Continuous distributions:

- Uniform: $X \sim U(a, b)$,
- Normal: $X \sim N(\mu, \sigma^2)$,
- “Standard” normal: $Z \sim N(0, 1)$,
- Exponential: $X \sim Exp(\lambda)$,
- Chi-square: $X \sim \chi_k^2$,
- Student's t: $X \sim t_\nu$,
- F-Snedecor: $X \sim F_{d_1, d_2}$.

Bernoulli distribution

PMF:

$$P(X = x) = p^x(1 - p)^{1-x}, \quad x \in \{0, 1\}, \quad 0 \leq p \leq 1$$





Example: a coin toss, success = heads, failure = tails.

Intuition: models an event with only two possible outcomes — success or failure.

Binomial distribution

PMF:

$$P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k}, \quad k = 0, 1, \dots, n$$





Example: number of heads in 10 coin tosses.

Intuition: the sum of n independent Bernoulli trials — how many successes occur in n trials.

Poisson distribution

PMF:

$$P(X = k) = \frac{\lambda^k e^{-\lambda}}{k!}, \quad k = 0, 1, 2, \dots$$





Example: number of phone calls received in one hour.

Intuition: models rare events occurring over a fixed time or space interval.

Uniform distribution

PDF:

$$f_X(x) = \frac{1}{b-a}, \quad x \in [a, b]$$





Example: drawing a random number from the interval $[0,1]$.

Intuition: all values within the interval are equally likely.

Normal (Gaussian) distribution

PDF:

$$f_X(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right), \quad x \in \mathbb{R}$$




Example: human height in a population.

Intuition: symmetric “bell-shaped” curve, with most values concentrated around the mean.

Exponential distribution

PDF:

$$f_X(x) = \lambda e^{-\lambda x}, \quad x \geq 0$$



Example: waiting time for a bus to arrive.



Intuition: models the time between events; small values are more probable.

Chi-square distribution

PDF:

$$f_X(x) = \frac{1}{2^{k/2}\Gamma(k/2)} x^{k/2-1} e^{-x/2}, \quad x \geq 0$$

$$\text{where } \Gamma(z) = \int_0^{\infty} t^{z-1} e^{-t} dt, \quad z > 0$$



Example: goodness-of-fit or variance tests in statistics.


Intuition: the sum of squares of k independent standard normal variables.

Student's t distribution

PDF:

$$f_X(x) = \frac{\Gamma(\frac{\nu+1}{2})}{\sqrt{\nu\pi}\Gamma(\frac{\nu}{2})} \left(1 + \frac{x^2}{\nu}\right)^{-\frac{\nu+1}{2}}, \quad x \in \mathbb{R}$$

where $\Gamma(z) = \int_0^\infty t^{z-1} e^{-t} dt, \quad z > 0$



Example: testing means with a small sample size.



Intuition: similar to the normal distribution, but with heavier tails for small

F-Snedecor distribution

PDF:

$$f_X(x) = \frac{\sqrt{\frac{(d_1x)^{d_1} d_2^{d_2}}{(d_1x+d_2)^{d_1+d_2}}}}{x B(d_1/2, d_2/2)}, \quad x \geq 0$$

$$\text{where } B(a, b) = \int_0^1 t^{a-1} (1-t)^{b-1} dt, \quad a, b > 0$$



Example: analysis of variance (ANOVA).

Intuition: compares the variances of two independent samples.

Relationships between distributions

Derivations of key distributions:

- Chi-square:

$$X_1, \dots, X_k \sim N(0, 1) \Rightarrow \chi_k^2 = \sum_{i=1}^k X_i^2$$

- Student's t:

$$Z \sim N(0, 1), \quad Y \sim \chi_\nu^2 \text{ independent} \Rightarrow T = \frac{Z}{\sqrt{Y/\nu}} \sim t_\nu$$

- F-Snedecor:

$$X \sim \chi_{d_1}^2, \quad Y \sim \chi_{d_2}^2 \text{ independent} \Rightarrow F = \frac{X/d_1}{Y/d_2} \sim F_{d_1, d_2}$$

Intuition:

- Chi-square = sum of squared standard normal variables,
- Student's t = standard normal divided by the square root of a scaled Chi-square variable,
- F-Snedecor = ratio of two scaled Chi-square variables.

Statistical hypothesis testing

What is statistical hypothesis testing?

- Hypothesis testing is a statistical method used to make inferences about a **population** based on **sample** data.
- It allows us to decide whether there is **enough evidence to reject a claim** (hypothesis) about a population parameter.
- There is a formal framework to do the above; we consider two competing hypotheses in it:

H_0 : null hypothesis (status quo)

H_1 : alternative hypothesis (research hypothesis)

Type I and type II errors

	fail to reject H_0	reject H_0
H_0 is true	correct decision	type I Error (α)
H_0 is false	type II Error (β)	correct decision

The most important ideas:

- **type I error** - rejecting H_0 when it is true,
- **type II error** - failing to reject H_0 when it is false,
- **significance level (α)**: probability of type I error, it is usually assumed (i.e., we assume what type I error we can accept when rejecting true H_0), common choices are: 0.05, 0.01, or 0.10,
- **power of the test ($1 - \beta$)**: probability of correctly rejecting H_0 when H_1 is true.

The hypothesis testing procedure

- ① Formulate H_0 and H_1 .
- ② Choose a significance level α .
- ③ Select an appropriate test statistic.
- ④ Determine the rejection region or critical value.
- ⑤ Compute the test statistic from the sample.
- ⑥ Make a decision: reject or fail to reject H_0 .
- ⑦ Interpret the result in context.

Example: one-sample t-test

Considered scenario:

- you have one group of numbers (e.g., test scores, heights, etc.),
- you want to check if the mean of your sample is different from a specific number.

Assumptions:

- continuous data: the variable is measured on a continuous scale,
- independence: each observation is independent from the others,
- normality: the population (or the sample) should be approximately normally distributed.

Numeric example:

- sample data: [102, 98, 101, 105, 97, 99, 100, 103, 95, 104],
- population mean: 100,
- question: is the sample mean different from 100?

Statistical framework:

- null hypothesis (H_0): $\mu = 100$ (no difference),
- alternative hypothesis (H_1): $\mu \neq 100$ (there is a difference).

Example: one-sample t-test - cont'd

Test statistic:

$$t = \frac{\bar{X} - \mu_0}{s/\sqrt{n}} \sim t_{n-1}$$

Inference:

$$|t| > t_{\alpha/2, n-1} \Rightarrow \text{reject } H_0 \text{ in favour of } H_1$$

Intuition:

- the test compares the **difference between your sample mean and 100** relative to the **spread of your data**, i.e., the test statistic shows how large the difference between the sample mean and 100 is (in standard-error units),
- if the difference is big compared to the spread \rightarrow not likely by chance, if it's small or your data are noisy \rightarrow could be just random.

Example: one-sample t-test - cont'd

Numeric example:


- sample data: [102, 98, 101, 105, 97, 99, 100, 103, 95, 104],
- population mean: 100,
- $n = 10$, $\bar{X} = 100.4$, $s = 3.03$, $\mu_0 = 100$,
- test statistic computation:

$$t = \frac{100.4 - 100}{3.03/\sqrt{10}} \approx 0.42$$

- critical value, assuming $\alpha = 0.05$:

$$t_{0.025,9} = 2.262 \Rightarrow |t| < t_{0.025,9} \Rightarrow \text{fail to reject } H_0$$

Example: one-sample t-test - cont'd



Example: two-sample t-test

Considered scenario:

- you have two independent groups (e.g., men vs women, treatment vs control),
- you want to check if their means are significantly different.

Assumptions:

- continuous data,
- independence: observations in each group are independent,
- normality: both groups come from approximately normal populations,
- equal variances (for the standard version; if not, use Welch's test).

Numeric example:

- group A (control): [102, 98, 101, 105, 97],
- group B (treatment): [110, 108, 112, 107, 111],
- question: is there a significant difference between group means?

Example: two-sample t-test - cont'd

Statistical framework:

- null hypothesis (H_0): $\mu_1 = \mu_2$ (no difference),
- alternative hypothesis (H_1): $\mu_1 \neq \mu_2$ (difference exists).

Test statistic (equal variances):

$$t = \frac{\bar{X}_1 - \bar{X}_2}{s_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}, \quad s_p = \sqrt{\frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2}{n_1 + n_2 - 2}}$$
$$t \sim t_{n_1+n_2-2}$$

Inference:

$$|t| > t_{\alpha/2, n_1+n_2-2} \Rightarrow \text{reject } H_0$$

Intuition:

- compares the difference between two sample means relative to the spread within groups,
- large $|t| \rightarrow$ unlikely that observed difference is due to chance,
- small $|t| \rightarrow$ the means are similar given the variability.

Example: two-sample t-test - cont'd

Numeric computation:

- $\bar{X}_1 = 100.6, s_1 = 3.2, n_1 = 5,$
- $\bar{X}_2 = 109.6, s_2 = 1.9, n_2 = 5,$
- pooled variance:

$$s_p = \sqrt{\frac{(4)(3.2^2) + (4)(1.9^2)}{8}} = 2.68$$


- test statistic:

$$t = \frac{100.6 - 109.6}{2.68 \sqrt{\frac{1}{5} + \frac{1}{5}}} = -6.0$$

- critical value for $\alpha = 0.05, df = 8: t_{0.025,8} = 2.306$

$$|t| > t_{0.025,8} \Rightarrow \text{reject } H_0$$

Example: two-sample t-test - cont'd



Example: two-sample t-test (left-tailed)

Statistical framework:

- null hypothesis (H_0): $\mu_1 = \mu_2$,
- alternative hypothesis (H_1): $\mu_1 < \mu_2$.

Test statistic:

$$t = \frac{\bar{X}_1 - \bar{X}_2}{s_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}$$
$$t \sim t_{n_1+n_2-2}$$

Decision rule:

$$t < -t_{\alpha, n_1+n_2-2} \Rightarrow \text{reject } H_0$$

Example:

$$t = -6.0, \quad t_{0.05,8} = 1.860$$

$$t < -t_{0.05,8} \Rightarrow -6.0 < -1.860 \Rightarrow \text{reject } H_0$$

Conclusion: mean of group 1 is significantly smaller than mean of group 2.

Example: two-sample t-test (right-tailed)

Statistical framework:

- null hypothesis (H_0): $\mu_1 = \mu_2$,
- alternative hypothesis (H_1): $\mu_1 > \mu_2$.

Test statistic:

$$t = \frac{\bar{X}_1 - \bar{X}_2}{s_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}$$
$$t \sim t_{n_1+n_2-2}$$

Decision rule:

$$t > t_{\alpha, n_1+n_2-2} \Rightarrow \text{reject } H_0$$

Example:

$$t = -6.0, \quad t_{0.05,8} = 1.860$$

$$t > t_{0.05,8} \Rightarrow -6.0 > 1.860 \Rightarrow \text{do not reject } H_0$$

Conclusion: no evidence that mean of group 1 is significantly greater than mean of group 2.

What is the p-value?

Definition

The **p-value** is the probability of obtaining a result at least as extreme as the one observed, assuming that the null hypothesis (H_0) is true.

- Small p-value ($< \alpha$) \Rightarrow evidence **against** H_0 .
- Large p-value ($> \alpha$) \Rightarrow data are **consistent** with H_0 .

Note: α is a preassumed significance level, commonly $\alpha = 5\%$, sometimes $\alpha = 10\%$, or $\alpha = 1\%$, or $\alpha = 0.1\%$.

Interpreting p-values obtained from statistical tests 1/3

Example (revisited) - one-sample t-test:

sample = [102, 98, 101, 105, 97, 99, 100, 103, 95, 104]

$$n = 10, \quad \bar{X} = 100.4, \quad s = 3.03, \quad \mu_0 = 100$$

$$t = \frac{100.4 - 100}{3.03/\sqrt{10}} \approx 0.42$$

$$\text{degrees of freedom: } \nu = n - 1 = 9$$

Two-sided p-value:

$$p = 2(1 - F_{t,9}(|t|)) \approx 0.68$$

Conclusion:

$$p \approx 0.68 > 0.05 \Rightarrow \text{fail to reject } H_0$$

i.e., no evidence that the mean differs from 100

Interpreting p-values obtained from statistical tests 2/3

Example (revisited) - two-sample t-test:

group A = [102, 98, 101, 105, 97], group B = [110, 108, 112, 107, 111]

$$n_1 = n_2 = 5, \quad \bar{X}_1 = 100.6, \quad \bar{X}_2 = 109.6, \quad s_1 = 3.2, \quad s_2 = 1.9$$

$$s_p = \sqrt{\frac{(4)(3.2^2) + (4)(1.9^2)}{8}} \approx 2.68$$

$$t = \frac{100.6 - 109.6}{2.68\sqrt{1/5 + 1/5}} \approx -6.0$$

degrees of freedom: $\nu = n_1 + n_2 - 2 = 8$

Two-sided p-value:

$$p = 2(1 - F_{t,8}(|t|)) \approx 0.0002$$

Conclusion:

$$p \approx 0.0002 < 0.05 \Rightarrow \text{reject } H_0$$

i.e., strong evidence that the means of group A and B differ.

Interpreting p-values obtained from statistical tests 3/3

Example - two sample t-test, left-tailed:

$$H_0 : \mu_1 = \mu_2, \quad H_1 : \mu_1 < \mu_2$$

$$t \approx -6.0, \quad \nu = 8$$

Left-tailed p-value:

$$p = F_{t,8}(t) \approx 0.0001$$

Conclusion:

$$p \approx 0.0001 < 0.05 \Rightarrow \text{reject } H_0$$

i.e., mean of group A is significantly smaller than mean of group B.

Thank you for your attention!

Maciej Świtała, PhD

ms.switala@uw.edu.pl